

HW1: Optimizing Matrix Multiply Report

Lian Liao, ll987, ll987(Perlmutter)

Results:

Description: Lian's simple blocked dgemm.			
Size: 31	Mflops/s: 18215.76	Percentage: 32.53	
Size: 32	Mflops/s: 21380.83	Percentage: 38.18	
Size: 96	Mflops/s: 24341.68	Percentage: 43.47	
Size: 97	Mflops/s: 22339.48	Percentage: 39.89	
Size: 127	Mflops/s: 21275.31	Percentage: 37.99	
Size: 128	Mflops/s: 24527.16	Percentage: 43.80	
Size: 129	Mflops/s: 21952.85	Percentage: 39.20	
Size: 191	Mflops/s: 22069.91	Percentage: 39.41	
Size: 192	Mflops/s: 24809.92	Percentage: 44.30	
Size: 229	Mflops/s: 23531.40	Percentage: 42.02	
Size: 255	Mflops/s: 22415.87	Percentage: 40.03	
Size: 256	Mflops/s: 23466.10	Percentage: 41.90	
Size: 257	Mflops/s: 22671.96	Percentage: 40.49	
Size: 319	Mflops/s: 23010.51	Percentage: 41.09	
Size: 320	Mflops/s: 24729.34	Percentage: 44.16	
Size: 321	Mflops/s: 23424.00	Percentage: 41.83	
Size: 417	Mflops/s: 24012.34	Percentage: 42.88	
Size: 479	Mflops/s: 22569.95	Percentage: 40.30	
Size: 480	Mflops/s: 24623.90	Percentage: 43.97	
Size: 511	Mflops/s: 22161.70	Percentage: 39.57	
Size: 512	Mflops/s: 21316.95	Percentage: 38.07	
Size: 639	Mflops/s: 23728.70	Percentage: 42.37	
Size: 640	Mflops/s: 24337.85	Percentage: 43.46	
Size: 767	Mflops/s: 23745.19	Percentage: 42.40	
Size: 768	Mflops/s: 23453.67	Percentage: 41.88	
Size: 769	Mflops/s: 23736.84	Percentage: 42.39	
Average percentage of Peak = 41.06			

Description: Lian's simple blocked dgemm.			
Size: 31	Mflops/s: 19580.72	Percentage: 34.97	
Size: 32	Mflops/s: 21468.34	Percentage: 38.34	
Size: 96	Mflops/s: 17899.13	Percentage: 31.96	
Size: 97	Mflops/s: 14008.13	Percentage: 25.01	
Size: 127	Mflops/s: 16502.12	Percentage: 29.47	
Size: 128	Mflops/s: 16858.31	Percentage: 30.10	
Size: 129	Mflops/s: 14870.02	Percentage: 26.55	
Size: 191	Mflops/s: 18898.74	Percentage: 33.75	
Size: 192	Mflops/s: 19207.89	Percentage: 34.30	
Size: 229	Mflops/s: 19511.36	Percentage: 34.84	
Size: 255	Mflops/s: 16171.19	Percentage: 28.88	
Size: 256	Mflops/s: 16349.90	Percentage: 29.20	
Size: 257	Mflops/s: 19309.90	Percentage: 34.48	
Size: 319	Mflops/s: 21811.90	Percentage: 38.95	
Size: 320	Mflops/s: 22010.27	Percentage: 39.30	
Size: 321	Mflops/s: 20843.82	Percentage: 37.22	
Size: 417	Mflops/s: 22353.49	Percentage: 39.92	
Size: 479	Mflops/s: 23821.97	Percentage: 42.54	
Size: 480	Mflops/s: 23995.03	Percentage: 42.85	
Size: 511	Mflops/s: 16456.49	Percentage: 29.39	
Size: 512	Mflops/s: 16590.15	Percentage: 29.63	
Size: 639	Mflops/s: 23424.52	Percentage: 41.83	
Size: 640	Mflops/s: 23564.24	Percentage: 42.08	
Size: 767	Mflops/s: 20161.97	Percentage: 36.00	
Size: 768	Mflops/s: 20272.97	Percentage: 36.20	
Size: 769	Mflops/s: 28075.46	Percentage: 50.13	
Average percentage of Peak = 35.30			

My best try is around 41% of the peak. To get this result, I have tried two main approaches. The main difference is that one initially realigns and fixes the size to memory; the other one doesn't realign it and keeps the size then deals with a tail. For me, the approach of dealing with the tail is better, which is 40%. And I combine two of them to get the 41%

optimizations attempt:

Based on the HW1 doc. I tried to implement multi-level blocking. I just simply added one more layer with 3 for-loop.

```

// For each block-row of A
for (int i = 0; i < fixlda; i += BLOCK_SIZE) {
    // For each block-row of A
    for (int j = 0; j < fixlda; j += BLOCK_SIZE) {
        // For each block-column of B
        for (int k = 0; k < fixlda; k += BLOCK_SIZE) {
            int M = min(BLOCK_SIZE, fixlda - i);
            int N = min(BLOCK_SIZE, fixlda - j);
            int K = min(BLOCK_SIZE, fixlda - k);
            // do_block(fixlda, M, N, K, A_block + i + k * fixlda,
            // Perform individual block dgemm
            for (int f = i; f < i + M; f += SMALL_BLOCK) {
                for (int g = j; g < j + N; g += SMALL_BLOCK) {
                    for (int h = k; h < k + K; h += SMALL_BLOCK) {
                        int Q = min(SMALL_BLOCK, i + M - f);
                        int W = min(SMALL_BLOCK, j + N - g);
                        int E = min(SMALL_BLOCK, k + K - h);

```

Use the same logic to add a second layer. The result is positive, but not increase speed much from just one level of blocking.

Next one is Repack and Realign

```

double* A_block = (double*)_mm_malloc(size:M_mod_4 * K * sizeof(double), align:32);
double* B_block = (double*)_mm_malloc(size:N_mod_4 * K * sizeof(double), align:32);

```

I set 32 bytes because my avx is 256. $256/32=8$ which is double's size.

I didn't do the Repack. I think it can potentially increase the speed with blocking.

```

for (j = 0; j < N - 3; j += 4) {
    b_ptr = &B_block[j * K];

    for (int m = 0; m < K; m++) {
        for (int n = 0; n < 4; n++) {
            b_ptr[m * 4 + n] = B[(j + n) * lda + m];
        }
    }

    for (i = 0; i < M - 3; i += 4) {
        a_ptr = &A_block[i * K];

        if (j == 0) {
            double* a_src = A + i;
            for (int u = 0; u < K; u++) {
                memcpy(dest:a_ptr + u * 4, src:a_src, n:4 * sizeof(double));
                a_src += lda;
            }
        }

        c_ptr = C + i + j * lda;
        kernel(lda, K, A:a_ptr, B:b_ptr, C:c_ptr);
    }
}

```

Here is in small block(C is M-by-N, A is M-by-K, and B is K-by-N)

I copied the matrix within the loop. Step is 4 because memory size is 32 and double is 8, $32/8=4$. I only do Realign for A and B because it is slower to add C. I have tried to implement C_block in small blocks or initials, but they are slower. I guess because they are not really aligned since I didn't fix the size before getting the microkernel.

My micro-kernel is inspired by doc

```
void micro_kernel (double* A, double* B, double* C) {
    // Declare
    __m512d Ar;
    __m512d Br;
    __m512d Cr;

    // Load
    Ar = _mm512_load_pd(A);
    Br = _mm512_load_pd(B);

    // Compute
    Cr = _mm512_add_pd(Ar, Br);

    // Store
    _mm512_store_pd(C, Cr);
}
```

But I use __m256d is not supported in Perlmutter. My loop for K is unrolling for every

```
for (int i = 0; i < KK; i += 2){
    aa1 = _mm256_load_pd(p:A);
    A += 4;

    bb01 = _mm256_broadcast_sd(a:B++);
    bb02 = _mm256_broadcast_sd(a:B++);
    bb03 = _mm256_broadcast_sd(a:B++);
    bb04 = _mm256_broadcast_sd(a:B++);

    cc01 = _mm256_fmadd_pd(A:aa1, B:bb01, C:cc01);
    cc02 = _mm256_fmadd_pd(A:aa1, B:bb02, C:cc02);
    cc03 = _mm256_fmadd_pd(A:aa1, B:bb03, C:cc03);
    cc04 = _mm256_fmadd_pd(A:aa1, B:bb04, C:cc04);

    aa2 = _mm256_load_pd(p:A);
    A += 4;

    bb11 = _mm256_broadcast_sd(a:B++);
    bb12 = _mm256_broadcast_sd(a:B++);
    bb13 = _mm256_broadcast_sd(a:B++);
    bb14 = _mm256_broadcast_sd(a:B++);

    cc11 = _mm256_fmadd_pd(A:aa2, B:bb11, C:cc11);
    cc12 = _mm256_fmadd_pd(A:aa2, B:bb12, C:cc12);
    cc13 = _mm256_fmadd_pd(A:aa2, B:bb13, C:cc13);
    cc14 = _mm256_fmadd_pd(A:aa2, B:bb14, C:cc14);
}
```

2 steps.

I added them while in the loop and combined them in the end, then wrote to C. After implementing the Micro-kernel, the speed increased a lot. This verified the results from some HPC papers which blocking and micro-kernel are most effective.

```
if (M%4 != 0){
    for (; i < M; ++i){
        for (int p = 0; p < N; p++){
            tem = C[i + p * lda];
            for (int k = 0; k < K; k++){
                tem += A[i + k * lda] * B[k + p * lda];
            }
            C[i + p * lda] = tem;
        }
    }
}
if (N%4 != 0){
    for (; j < N; j++){
        for (int p = 0; p < M_mod_4; p++){
            tem = C[p + j * lda];
            for (int k = 0; k < K; k++){
                tem += A[p + k * lda] * B[k + j * lda];
            }
            C[p + j * lda] = tem;
        }
    }
}
```

For the tail. Just simply access the original A,B, and C to deal. I tried to make them aligned to avoid tail, but it didn't work out.

Other small optimizations:

```
#pragma GCC optimize ("peel-loops")
#pragma GCC optimize("inline")
#pragma GCC optimize("unroll-loops")
#pragma GCC optimize("Ofast")
```

Those are optimizers I found. They can help my local computer dramatically speed up, but not for Perlmutter. The interesting thing is that with optimizers, my computer is faster than in Perlmutter; without them, it is way slower than in Perlmutter.

I read that an "inline" function can help with optimization because the compiler attempts to embed the function's code directly into the calling site, rather than generating a function call. However, I don't see any difference after adding them to


my function.

Same as “restrict” I don’t see any difference after adding them to my Matrix.

```
#ifndef BLOCK_SIZE
#define BLOCK_SIZE 128
#define SMALL_BLOCK 64
#endif
```

This is the critical factor for performance. Is decided whether the blocking size or represented as L1 and L2 cache. (128,64) is my best set to my code.

Note: because my other approach is better on size 31. So I only run it if the size is smaller than 32



This is my plot for the final performance. My curve is flat at around 40%

Conclusion:

The optimization attempts on matrix multiplication yielded a significant performance increase, achieving approximately 41% of the peak. The critical factor for performance optimization was the cache-aware blocking size, indicating the profound impact of memory hierarchy on computational efficiency (speed would drop a lot if I

forgot to free it). Despite the various strategies employed, there remains a performance gap to the theoretical peak.

References

Jiang, Xuan, et al. "Optimizing Matrix Multiplication on Nersc's High Performance Computer Cori." OSF Preprints, 26 Feb. 2022. Web.